



# Digital Techniques 2

## L5: Digital Arithmetic



# Fixed-Point Representation of Numbers

- Fixed-point (FXP) numbers have a fixed number of digits on the right hand side of the radix point. e.g.

$$0010_2 = 2_{10}$$

$$00.10_2 = 0.5_{10}$$

$$01.01_2 = 1.25_{10}$$

- The radix point is not represented in hardware, so you have to keep track of it!
- Signed FXP numbers are usually encoded in two'2 complement format



Two's complement encoding method:

1. Select the number of bits in codewords (N)
2. Encode numbers  $X \in [-2^{N-1}, 2^{N-1}-1]$  as  $2^N + X$  and **drop the N+1:th bit**, using codes [10..00, 10..01, ..., 01..11]
3. Example: N = 4 bits,  
 $X = -1$   
 $X_{2^4's\ complement} = 2^4 + (-1)$   
 $= 16 - 1 = 10000 - 0001 = 01111$
4. We can now add positive and negative numbers:  
 $2 + (-1) = 0010 + 1111 = 10001$

To do: Check that this works for all positive and negative numbers. Study 2's complement number encoding until you understand how it works!



# Floating-Point Representation of Number

- Floating point number representation:  
**significant \* base <sup>exponent</sup>**
- In digital circuits, significant (S) and exponent (E) are (signed) fixed-point numbers, and base = 2
- FLP numbers are usually encoded by...
  - S = sign bit and unsigned value
  - Normalization: leading zeros in S value are omitted and E is adjusted accordingly
  - Exponent biasing: a constant value is added to exponent to make it positive (sign bit does not have to be stored)
- Floating point numbers have a (much) **larger dynamic range** [min, max] but **lower precision** than fixed point numbers
- In SystemVerilog the *float* type is **real**
- In general, floating point hardware is more complex than fixed-point
- Synthesis tools in general do not support floating point numbers, but that may be changing
- **In this course, we only use fixed-point numbers**

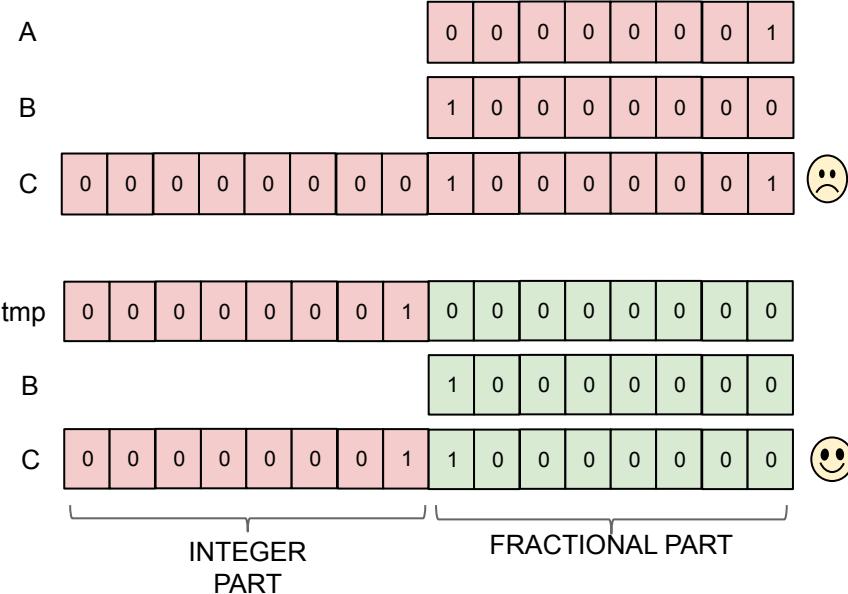


# Working with Fixed-Point Numbers

- You have to take care that radix points stay aligned in operations
- Example: Assume A has radix point at right end, and B at left end

```
logic [7:0] A, B;  
logic [15:0] C, tmp;  
A = 8'b00000001;  
B = 8'b10000000;  
  
C = A + B;    ☹  
  
tmp = A << 8;  
C = tmp + B  ☺
```

Let A = 00000001. ( $1_{10}$ ) and B = .10000000 ( $0.5_{10}$ )

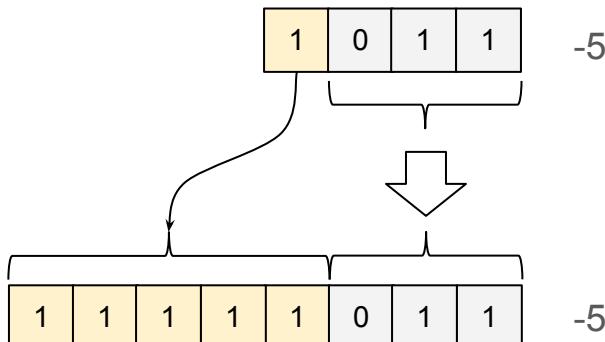


**Tip:** It's a good idea to use "temp" variables to make sure sign and bit-length are handled correctly,



# Manipulation of Signed FXP Numbers

- To obtain the **complement** of a 2's complement number, first invert all bits and then add a '1' to the inverted value:  
 $-111101 = 000010 + 1 = 000011 = 3$
- To make a signed FXP number longer, **sign extension** must be done:



- Shortening a number from left end:  
Value may be outside the new range, is often **saturated** (clipped to max possible value) before shortening.  
Example: Shorten 01111, 10001 and 00111 to four bits (to range [-8, 7])  
 $01111 \text{ ( } 15_{10} \text{) } \Rightarrow 0111 \text{ ( } 7_{10} \text{) }$   
 $10001 \text{ ( } -15_{10} \text{) } \Rightarrow 1000 \text{ ( } -8_{10} \text{) }$   
 $00111 \text{ ( } 7_{10} \text{) } \Rightarrow 0111 \text{ ( } 7_{10} \text{) }$
- Shortening from the right end:
  - **Truncation** (just drop bits from the end)
  - **Rounding + truncation** (better, but requires comparator hardware)



# Handling of Signed FXP Numbers in SystemVerilog Designs

- **logic** vector works like an unsigned binary number by default, but you can change the encoding if necessary:

```
logic      [7:0] a;
logic unsigned [7:0] b;
logic signed   [7:0] c;
```
- SystemVerilog is a *weakly typed* language: operations on data items that are of incompatible types are allowed **according to language rules**
- Signedness handling is an area where you certainly have to know the rules

Assume a = 1110 (14)

```
module confused (input logic [3:0] a,
                  output logic less_than_1,
                  output logic less_minus_1);
    always_comb
        begin
            if (a < 1)
                less_than_1= '1';
            else
                less_than_1= '0';

            if (a < -1)
                less_minus_1= '1';
            else
                less_minus_1= '0';
        end
    endmodule
```

(14 < 1)

(14 < 15)

SV Standard 11.4.4, "When one or both operands of a relational expression are unsigned, the expression shall be interpreted as a comparison between unsigned values."



# Some of SystemVerilog's Signedness Rules

- Decimal number literals (such as -1 or 4) are signed
- Based literal constants (`4'b1111`) are unsigned unless indicated with an 's' (`4'sb1111`)
- Packed logic arrays (`logic [3:0] x`) are unsigned unless declared *signed*
- Bit selects (`y[3]`) and range selects (`y[3:0]`) are always unsigned (gotcha!)
- Concatenation `{y, x}` result is unsigned
- Sign of an assignment depends only on operands ( $z = x + y$ ), not on left side
- An expression is signed only if all operands are signed

How to be sure? Use `$signed` and `$unsigned` conversion functions, e.g.

instead of `if (a < -1)` write `if ($signed(a) < -1)` if a should be treated as signed

Survival guide: If you are uncertain, write a simple testbench program and try it out! (like in `labs/lab_datatypes`)



# Some of SystemVerilog's Bit-Length Rules

- int and unspecified-length literals: 32 bits
- Operation  $i \text{ op } j$  where  $\text{op}$  is  $+, -, *, /, \%, ^$  etc, ( $L(x)$  = length of  $x$ )  
 $L(\text{op}) = \max(L(i), L(j))$
- Operation  $i \text{ op } j$ , where  $\text{op}$  is  $<<, >>, **, <<<, \text{ or } >>>$   
 $L(\text{op}) = L(i)$
- In assignments, right-hand side length also depends on left hand side length

```
logic [3:0] p, q;  
logic [3:0] r;  
p = 7;  
q = 9;
```

```
r = (p + q) >> 1; // r == 0
```

$L = \max(4, 4) \Rightarrow 7+9 \text{ overflows}$

```
logic [3:0] p, q;  
logic [4:0] r;  
p = 7;  
q = 9;
```

```
r = (p + q) >> 1; // r == 8
```

$L(r) = 5 \Rightarrow \text{no overflow}$

```
logic [3:0] p, q;  
logic [3:0] r;  
p = 7;  
q = 9;
```

```
r = (7 + q) >> 1; // r == 8
```

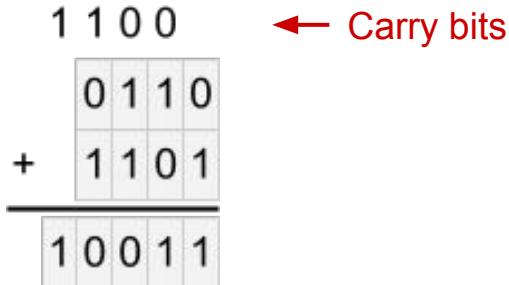
$L(7) = 32 \Rightarrow \text{no overflow}$



# Arithmetic Hardware: Addition

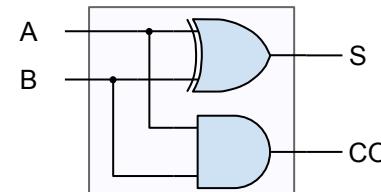
- Binary addition works just like they taught you in school:

$$A = 0110, B = 1101, A+B = 10011$$

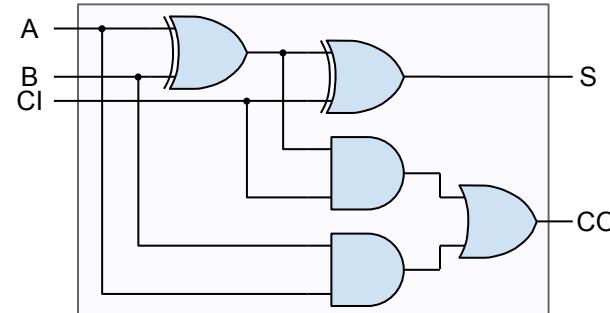
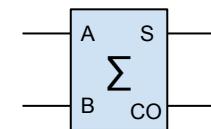


- Number of bits in sum  $A + B$  is **max( bits in A, bits in B ) + 1**
- To compute  $A - B$ , first form a 2's complement of B, then add

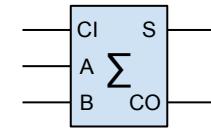
Building blocks:



Half adder



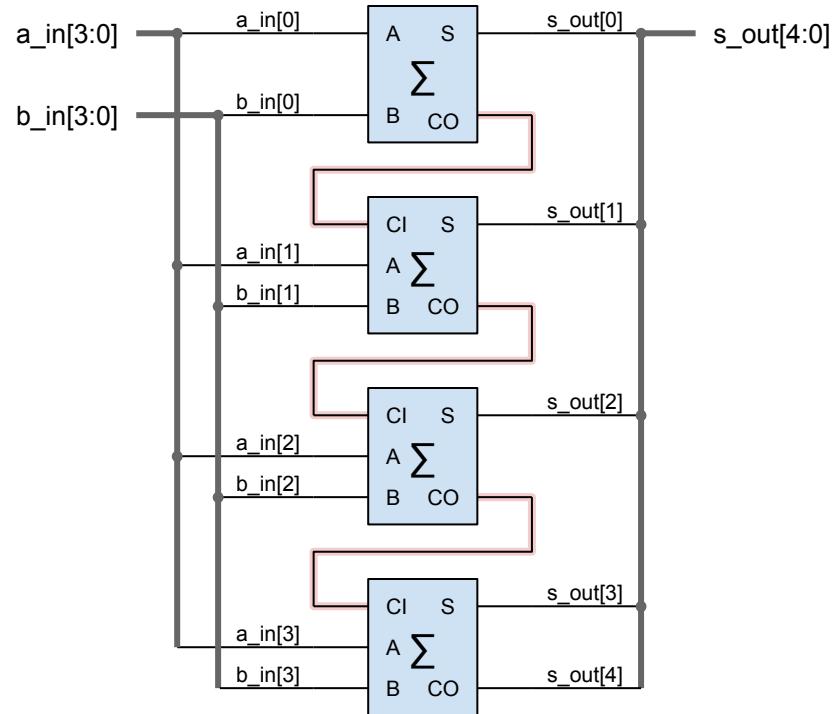
Full adder





# Ripple-Carry Adder

- Simplest binary adder circuit
- 'Ripple' refers to the basic feature and limitation: the carry bit **ripples** from adder stage to the next
- Also known as a **carry-propagate adder** (see also: [adder/subtractor](#))
- Delay and complexity are proportional to the number of bits in the addends
- Speeding up carry propagation is the main objective in developing faster adder architectures (at the cost of additional hardware)



Area: 17 gates  
Delay: 7 gates



# Carry Look-Ahead Adder

$g[i]$  = Carry is generated from bits  $a_{in}[i]$  and  $b_{in}[i]$  if

$$g[i] = a_{in}[i] \cdot b_{in}[i] \quad (\text{both are } 1)$$

$p[i]$  = Bits  $a_{in}[i]$  and  $b_{in}[i]$  propagate a 'lower' carry if

$$p[i] = a_{in}[i] \oplus b_{in}[i] \quad (\text{one is } 1)$$

The  $i$ :th carry bit then has the logic function:

$$c[i] = g[i] + p[i] \cdot c[i-1]$$

Three first carry bits of an adder:

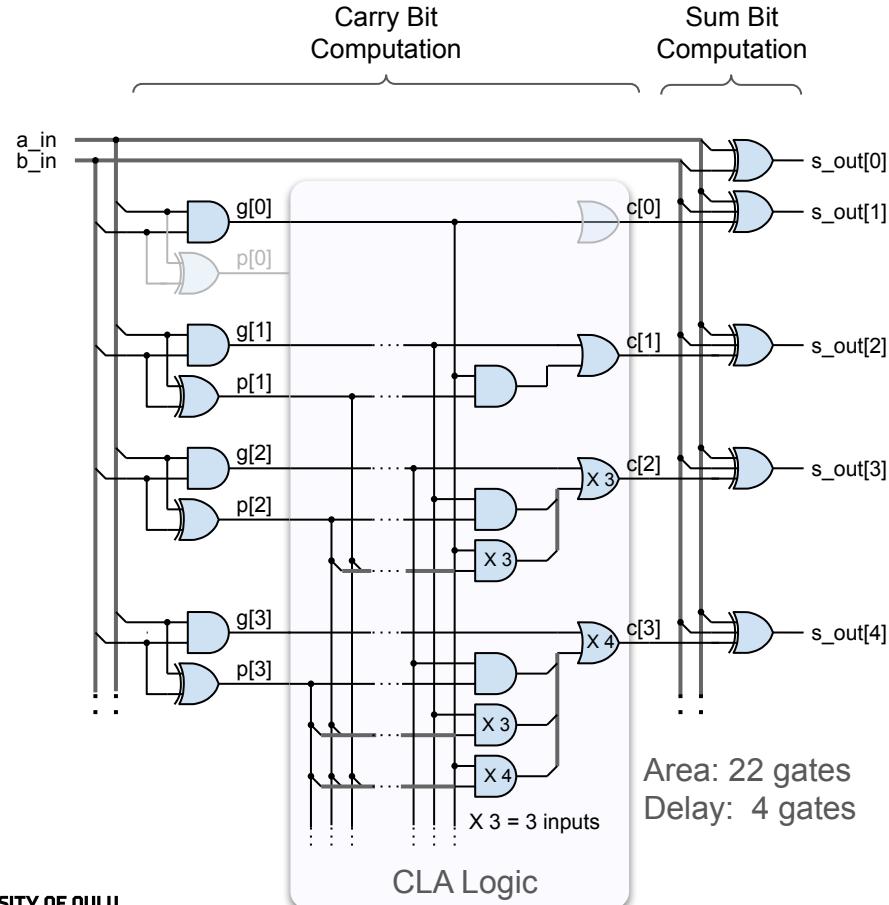
$$c[0] = g[0]$$

$$c[1] = g[1] + p[1] \cdot g[0]$$

$$c[2] = g[2] + p[2] \cdot g[1] + p[2] \cdot p[1] \cdot g[0]$$

Sum bit  $s_{out}[i]$  is given by the XOR function:

$$s_{out}[i] = a_{in}[i] \oplus b_{in}[1] \oplus c[i-1]$$

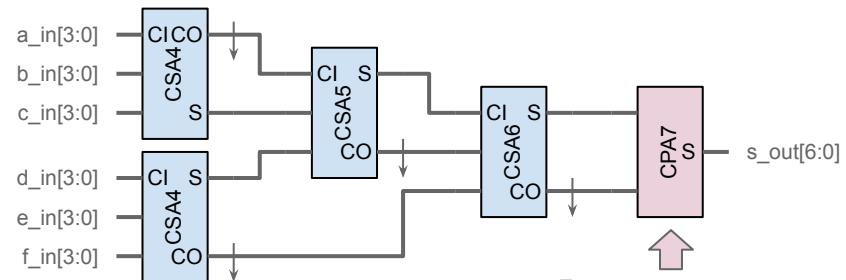
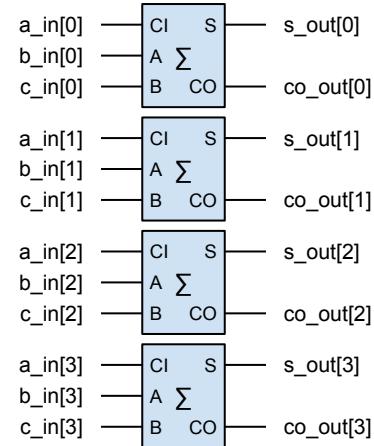




# Carry-Save Addition

- Observation:  $A+B =$   
 $(\text{Sum bits of } A+B) + (\text{Carry bits of } A + B) \ll 1$
- Sum and carry bits can be generated with **carry save adder** (CSA) that does not propagate carry
- Final sum is generated by adding sum bits with left-shifted carry bits in a carry propagate adder (CPA)
- To compute a **sum of many numbers**  $A+B+C+\dots$  we can delay the propagation of carry by creating intermediate sum and carry vectors in CSA:s and using a CPA as a last stage
- Optimal CSA tree is called a **Wallace tree**

CSA:





# Arithmetic Hardware: Multiplication

- Works like in school, again, except that partial products ( $P_0 - P_3$ ) are obtained with the AND operation (multiplication by 0 or 1)
- Implementation:
  - Generate partial products  $P$  with AND gates
  - Add partial products in a Wallace tree CSA
- Signed multiplication:
  - Convert operands to unsigned and fix the sign of the result
  - Use special algorithms, such as, [Booth](#)
- Number of bits in  $A * B$  is  
**bits in A + bits in B**

Example:  $5 \times 12 =$

$$\begin{array}{r} 1100 \\ \times 0101 \\ \hline 1100 \\ 0000 \\ 1100 \\ + 0000 \\ \hline 00111100 \end{array}$$

(12) (5)  
P0  
P1  
P2  
P3  
60

## How about other arithmetic operations?

There are countless logic architectures for implementing arithmetic operations (+, -, \*, /, >, <, sin, cos,...) in combinational or sequential logic, in bit-parallel and bit-serial format, with binary and 2's complement numbers etc. Consult the vast **computer arithmetic** literature to learn more.



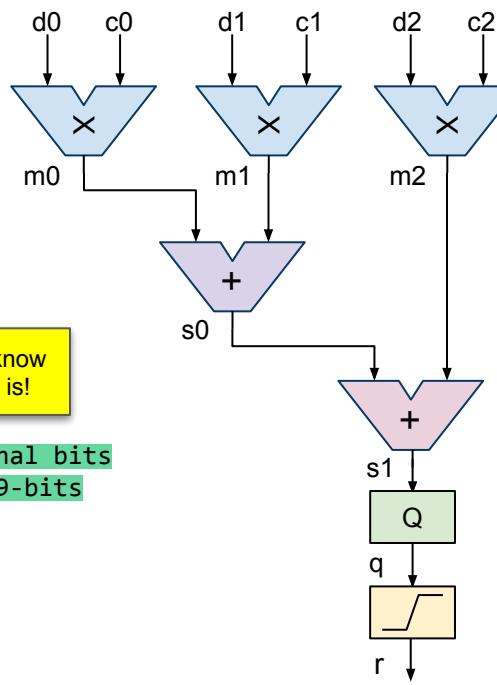
# Fixed-Point Datapath Sizing Example

```

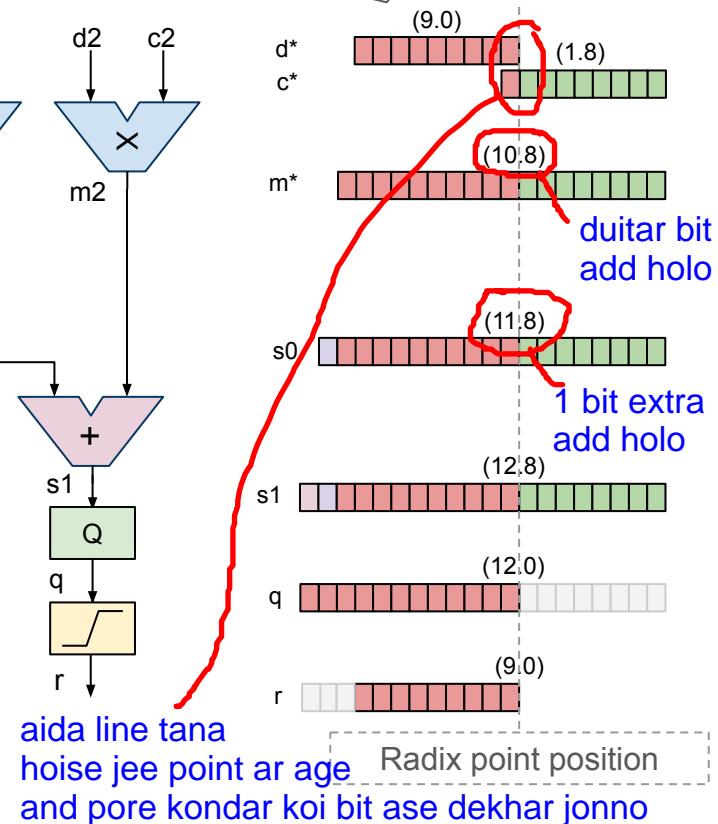
module ffp (input logic signed [8:0] c0, c1, c2,
            input logic signed [8:0] d0, d1, d2,
            output logic signed [8:0] r);
    logic signed [17:0]          m0, m1, m2;
    logic signed [18:0]          s0;
    logic signed [19:0]          s1;
    logic signed [11:0]          q;
    always_comb
        begin
            m0 = c0 * d0;
            m1 = c1 * d1;
            m2 = c2 * d2;
            s0 = m0 + m1;
            s1 = m2 + s0;
            q = s1 >>> 8;           // drop fractional bits
            if (q > 9'sb011111111)   // saturate to 9-bits
                r = 9'sb011111111;
            else if (q < 9'sb100000000)
                r = 9'sb100000000;
            else
                r = q;
        end
endmodule

```

Here we must know where the point is!



**Input and output formats:**  
d's and r: 9-bit signed integers  
c's: 9 bit FXP numbers with the point after MSB. Decim range is [-1.0, 0.99609375]

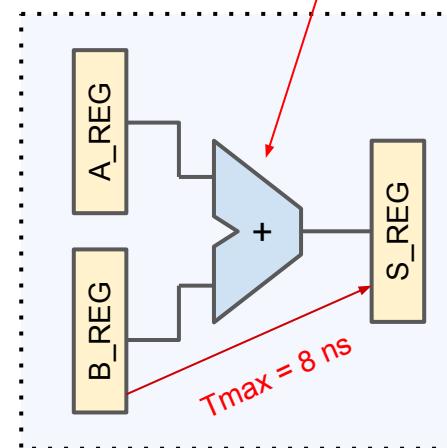




# Implementation of Arithmetic Operations in Synthesis

- Logic synthesis programs' RTL component libraries contain a selection of arithmetic components (+, -, \*, / etc)
- In *Design Compiler*, the RTL library is called *DesignWare*
- In RTL code translation, the synthesis program implements arithmetic operations detected in the code by inserting a library part in the unmapped design
- Library part implementation is automatically selected according to timing constraints in optimization
- Don't try to optimize the logic of an arithmetic part: it can make it slower (e.g. flattening to two-level will "optimize" CLA logic away!)

$S\_REG = A\_REG + B\_REG;$



RTL architecture generated  
from code.

RTL  
COMPONENT  
LIBRARY

CLA:  
delay=6  
area=11

RCA:  
delay=10  
area=5



# References

1. [IEEE Standard for SystemVerilog](#)
  - 11.6 Expression bit lengths
  - 11.8 Expression evaluation rules (signedness rules)
2. Harris (2007) Digital Design And Computer Architecture  
[Ch 5.2 - 5.3](#)